

Graded Homework 2, exercise 2

Marco Milanta

December 13, 2021

For this exercise I exchanged ideas with Jan Tempus.

Minimizing Cost (25 points)

Given a stream of n elements $x_1, \dots, x_n \in \{1, \dots, n\}$ all distinct (i.e., the input is a permutation). We have to choose x_i, x_j with $j \geq i$ on arrival, meaning that when an element arrives, we must immediately decide if we include it in one of the two elements. The cost function is given as $\frac{1}{x_j - x_i + 1}$. We would like to minimize this cost. (Note that any sensible strategy will have $\frac{1}{x_j - x_i + 1} < 1$ as we can choose $i = j = n$.)

1. Find a deterministic algorithm that achieves competitive ratio at most \sqrt{n} .
2. Show that any randomized algorithm has competitive ratio at least $\Omega(\sqrt{n})$

Solutions

1. We propose the following algorithm. Note that we index arrays starting from 1.

Legality of the algorithm. Before analyzing the competitive ratio of this algorithm, I would like to notice that it has the correct structure. Notice these two facts:

- In the k -th iteration of the loop, we only look at $X[k]$. This means that our algorithm doesn't cheat by looking into the future. Finally, we look at the last element after the loop.
- x_i and x_j are assigned only once, always x_i before x_j , and they are assigned to be $X[k]$ at the k -th iteration. Or, at latest, after the for loop. All of this is guaranteed by the counter `ChosenElements`. This means that we are respecting the constraint "when an element arrives, we must immediately decide if we include it in one of the two elements".
- The algorithm iterates over the input elements, therefore is polynomial in n .

Competitiveness of the algorithm. Now we want to analyze the performance. To do so, let $\hat{i}, \hat{j}, \hat{i} \leq \hat{j}$ be an optimal solution that minimize $\frac{1}{x_j - x_i + 1}$. We call OPT the **cost** of the optimal algorithm. Notice that the optimal cost cannot go lower $\frac{1}{n}$, this is achieved when we manage to get $x_i = 1, x_j = n$.

Theorem 1 (Algorithm 1 is \sqrt{n} -competitive). *The algorithm 1 is \sqrt{n} -competitive against the optimal solution. Or, formally, let i, j, cost be the output of the algorithm 1. Let OPT be the maximal profit, then, for any input*

$$\text{profit} \leq \sqrt{n}OPT.$$

Proof. We divide the proof in two parts depending on the behavior of the algorithm.

Algorithm 1: Online algorithm

Input: $X \in \text{permutations of } \{1, \dots, n\}$
Output: $i, j, \frac{1}{x_j - x_i + 1}$

```
1 Seen  $\leftarrow [0, \dots, 0] \in \mathbb{N}^n$ 
2 Missing  $\leftarrow [1, \dots, n] \in \mathbb{N}^n$ 
3 ChosenElements  $\leftarrow 0$ 
4  $x_i \leftarrow x_j \leftarrow 0$ 
5 for  $k \in 1, \dots, n-1$  do
6   if ChosenElements = 0 then
7     Missing[ $X[k]$ ]  $\leftarrow 0$ 
8     optFutu  $\leftarrow \max$  Missing
9     if (optFutu -  $X[k] + 1$ )  $\geq \sqrt{n}$  then
10       $x_i \leftarrow X[k]$ 
11       $i \leftarrow k$ 
12      ChosenElements  $\leftarrow 1$ 
13   else if ChosenElements = 1 then
14     if  $X[k] = \max$  Missing then
15        $x_j \leftarrow X[k]$ 
16        $j \leftarrow k$ 
17       ChosenElements  $\leftarrow 2$ 
18 if ChosenElements = 0 then
19    $x_i \leftarrow x_j \leftarrow X[n]$ 
20    $i \leftarrow j \leftarrow n$ 
21 return  $i, j, \frac{1}{x_j - x_i + 1}$ 
```

Case 1: (there is a $k \in 1, \dots, n-1$ for which we enter the **if** at line 9) Let for now k be the iteration in which we enter the **if** at line 9. One can notice that we will end up with a final cost of $\frac{1}{\text{optFutu} - X[k] + 1}$. This is because we chose $x_i = X[k]$, and then we chose x_j to be the maximal value among the remaining ones. We can safely find such value since we are keeping track of the missing ones (**optFutu**). Once **maxMissing** will arrive, we pick it for x_j . Notice that we cannot miss it since we are only looking at permutations as inputs. Hence, our algorithm has a cost of

$$\text{cost} \frac{1}{x_j - x_i + 1} = \frac{1}{\text{optFutu} - X[k] + 1}.$$

Using this, we get that:

$$\text{cost} = \frac{1}{\text{optFutu} - X[k] + 1} \stackrel{(i)}{\leq} \frac{1}{\sqrt{n}} = \sqrt{n} \frac{1}{n} \stackrel{(ii)}{\leq} \sqrt{n} OPT.$$

Where (i) follows from the fact that we enter the **if** at line 9, and therefore $\text{optFutu} - X[k] + 1 \geq \sqrt{n}$. (ii) follows from the fact that minimum possible cost is $\frac{1}{n}$. In this scenario we have shown that the algorithm is \sqrt{n} competitive.

Case 2: (we never enter the **if** at line 9) In this scenario, our algorithm picks $x_i = x_j = n$, and our **costs** will be 1. We now want to show that $OPT \geq \frac{1}{\sqrt{n}}$.

To do so, we continue by contradiction. Assume that the optimal solution takes $\hat{i} \leq \hat{j}$ such that $OPT = \frac{1}{x_{\hat{j}} - x_{\hat{i}} + 1} < \frac{1}{\sqrt{n}}$. This yields that

$$x_{\hat{j}} - x_{\hat{i}} + 1 > \sqrt{n}.$$

Now we look at the behavior of our algorithm in this setting. Let's observe the moment when it gets to step $k = \hat{i}$. In this step $\text{optFutu} = \max\text{Missing} \geq x_{\hat{j}}$ (this is true since $\hat{j} > \hat{i}$, it's not equal, otherwise OPT cost would be 1). This yields that $\text{optFutu} - X[\hat{i}] + 1 \geq x_{\hat{j}} - x_{\hat{j}} + 1 > \sqrt{n}$, which itself it yields that we enter the **if** at line 9. This is however a contradiction.

From this we have that $OPT \geq \frac{1}{\sqrt{n}}$, and we get:

$$\text{cost} = 1 = \sqrt{n} \frac{1}{\sqrt{n}} \leq \sqrt{n} OPT.$$

Conclusion: We have shown that independently of which part of the algorithm triggers, we can guarantee that it outputs a \sqrt{n} -competitive solution \square


2. Let $OPT(x)$ be the optimal **cost** achievable given $x \in \mathcal{X}$, where \mathcal{X} is the space of permutation of $\{1, \dots, n\}$. Let's define a distribution over inputs \mathcal{P} indexed by a discrete density $p : \mathcal{X} \rightarrow [0, 1]$:

$$p(x) = \begin{cases} \frac{1}{2} & x = \bar{x}_1 \\ \frac{1}{2} & x = \bar{x}_2 \\ 0 & \text{otherwise} \end{cases} \quad (1)$$

Where \bar{x}_1, \bar{x}_2 are just aliases for

$$\begin{aligned} \bar{x}_1 &= \{n-1, n-2, \dots, 1 + n - \sqrt{n}, n, n - \sqrt{n}, n - \sqrt{n} - 1, \dots, 1\} \\ \bar{x}_2 &= \{n-1, n-2, \dots, 1, n\}. \end{aligned}$$

Notice that we are assuming that \sqrt{n} is an integer (we don't generalize here, but we are sure that the idea generalizes as well). To visualize them better, look at figure below.



Where \bar{x}_1 is represented with blue dots and \bar{x}_2 with orange stars. In this plot $n = 64$.

Now we state a key lemma that shows the importance of the chosen \mathcal{P} .

Lemma 1. [Best distribution] Given $x \sim \mathcal{P}$, any possible deterministic algorithm \mathcal{A} has an expected competitive ratio of at least $\frac{1}{2}(\sqrt{n} + 1)$. Formally we have that,

$$\mathbb{E}_{x \sim \mathcal{P}} \left[\frac{\mathcal{A}(x)}{OPT(x)} \right] \geq \frac{1}{2}(\sqrt{n} + 1) \quad \text{for any } \mathcal{A} \text{ deterministic algorithm}$$

Proof. We start by defining two deterministic algorithms: \mathcal{A}_1 and \mathcal{A}_2 .

- \mathcal{A}_1 waits all the way to $1 + n - \sqrt{n}$, then buys $x_i = 1 + n - \sqrt{n}$. \mathcal{A}_1 then waits for n to appear and then picks it (note that it will always appear).
- \mathcal{A}_2 commits and wait all the way to the second to last element, and then picks it. \mathcal{A}_2 then picks the last element if n is the only missing element, and it buys the second to last again if the only missing element is 1.

\mathcal{A}_1 and \mathcal{A}_2 competitive ratio: looking at the behavior of the algorithms, and at how \bar{x}_1 and \bar{x}_2 are defined, we can compute the cost of the optimal algorithm, of \mathcal{A}_1 and of \mathcal{A}_2 on the inputs \bar{x}_1 and \bar{x}_2 :

$$\begin{aligned}
OPT(\bar{x}_1) &= \frac{1}{n-1-n-\sqrt{n}+1} = \frac{1}{\sqrt{n}} \\
OPT(\bar{x}_2) &= \frac{1}{n-1+1} = \frac{1}{n} \\
\mathcal{A}_1(\bar{x}_1) &= \frac{1}{n-1-n-\sqrt{n}+1} = \frac{1}{\sqrt{n}} \\
\mathcal{A}_1(\bar{x}_2) &= \frac{1}{n-1-n-\sqrt{n}+1} = \frac{1}{\sqrt{n}} \\
\mathcal{A}_2(\bar{x}_1) &= \frac{1}{n-1-(n-1)+1} = 1 \\
\mathcal{A}_2(\bar{x}_2) &= \frac{1}{n-1+1} = \frac{1}{n}.
\end{aligned}$$

Now we can compute:

$$\begin{aligned}
\mathbb{E}_{x \sim \mathcal{P}} \left[\frac{\mathcal{A}_1(x)}{OPT(x)} \right] &= p(\bar{x}_1) \frac{\mathcal{A}_1(\bar{x}_1)}{OPT(\bar{x}_1)} + p(\bar{x}_2) \frac{\mathcal{A}_1(\bar{x}_2)}{OPT(\bar{x}_2)} \\
&= \frac{1}{2} \left(\frac{\frac{1}{\sqrt{n}}}{\frac{1}{\sqrt{n}}} + \frac{\frac{1}{\sqrt{n}}}{\frac{1}{n}} \right) = \frac{1}{2} (1 + \sqrt{n})
\end{aligned}$$

And we also do it for \mathcal{A}_2 :

$$\begin{aligned}
\mathbb{E}_{x \sim \mathcal{P}} \left[\frac{\mathcal{A}_2(x)}{OPT(x)} \right] &= p(\bar{x}_1) \frac{\mathcal{A}_2(\bar{x}_1)}{OPT(\bar{x}_1)} + p(\bar{x}_2) \frac{\mathcal{A}_2(\bar{x}_2)}{OPT(\bar{x}_2)} \\
&= \frac{1}{2} \left(\frac{1}{\frac{1}{\sqrt{n}}} + \frac{\frac{1}{n}}{\frac{1}{n}} \right) = \frac{1}{2} (\sqrt{n} + 1)
\end{aligned}$$

Any other algorithm: Now we want to show that the result doesn't only hold for \mathcal{A}_1 and \mathcal{A}_2 , but also for any other deterministic algorithm. For this we group all the possible algorithms in 2 groups:

- **Group 1:** algorithms that pick as first element an element before $1 + n - \sqrt{n}$ both in \bar{x}_1 and \bar{x}_2
- **Group 2:** Every algorithm which is not in **Group 1**, and it's not \mathcal{A}_1 or \mathcal{A}_2 .

First we show that any algorithm in **Group 1** costs more than \mathcal{A}_1 independently of the input \bar{x}_1 and \bar{x}_2 . Let \mathcal{C} be an algorithm in **Group 1**: it picks the first element c_1 before $1 + n - \sqrt{n}$, and then it picks c_2 . This implies

$$\mathcal{C}(\bar{x}_1) = \mathcal{C}(\bar{x}_2) = \frac{1}{c_2 - c_1 + 1} \stackrel{(i)}{<} \frac{1}{c_2 - (1 + n - \sqrt{n}) + 1} \stackrel{(ii)}{\leq} \frac{1}{n - 1 - n + \sqrt{n} + 1} = \frac{1}{\sqrt{n}}$$

Where the inequality (i) holds since we c_1 is before $1 + n - \sqrt{n}$, and therefore $c_1 > 1 + n - \sqrt{n}$. (ii) holds since $c_2 \leq n$. Notice that both in \bar{x}_1 and \bar{x}_2 the cost of \mathcal{C} is more than $\frac{1}{\sqrt{n}} = \mathcal{A}_1(\bar{x}_1) = \mathcal{A}_1(\bar{x}_2)$. Therefore:

$$\mathbb{E}_{x \sim \mathcal{P}} \left[\frac{\mathcal{C}(x)}{OPT(x)} \right] \geq \mathbb{E}_{x \sim \mathcal{P}} \left[\frac{\mathcal{A}_1(x)}{OPT(x)} \right] \geq \frac{1}{2}(\sqrt{n} + 1).$$

This concludes the proof for **Group 1**.

We now notice that all the algorithms in **Group 2** buy the first element after $1 + n - \sqrt{n}$. This is not trivial since in **Group 1** there are the algorithms that pick the first element before $1 + n - \sqrt{n}$ both in \bar{x}_1 and \bar{x}_2 . What about the algorithms that depending on the input chose to pick either before or after $1 + n - \sqrt{n}$? Such algorithm cannot exist among deterministic ones, this is because the first elements before $1 + n - \sqrt{n}$ are exactly the same in \bar{x}_1 and \bar{x}_2 , therefore, a deterministic algorithm that picks an element before $1 + n - \sqrt{n}$ will do it in both inputs, and vice versa, if it doesn't pick it before $1 + n - \sqrt{n}$ it will not do it in both inputs.

Now, once we observed this, it is easy to notice that all algorithms in group to get a cost of at least 1 on the input \bar{x}_1 , this is because they "missed" the spike at $1 + n - \sqrt{n}$.

Now let \mathcal{B} be our algorithm in **Group 2**. Let c_1 be the element \mathcal{B} picks first on the input \bar{x}_2 , and c_2 the other. $c_1 > 1$, otherwise the algorithm is \mathcal{A}_2 , but **Group 2** doesn't contain \mathcal{A}_2 . And, as before, $c_2 \leq n$. This yields that the cost in input \bar{x}_2 is

$$\mathcal{B}(\bar{x}_2) = \frac{1}{c_2 - c_1 + 1} \geq \frac{1}{n - c_1 + 1} > \frac{1}{n - 1 + 1} = \frac{1}{n}.$$

Furthermore, from the observation before, we have $\mathcal{B}(\bar{x}_1) \geq 1$. Finally, we can observe that \mathcal{B} does strictly worse than \mathcal{A}_2 both in \bar{x}_1 and \bar{x}_2 . Therefore

$$\mathbb{E}_{x \sim \mathcal{P}} \left[\frac{\mathcal{B}(x)}{OPT(x)} \right] \geq \mathbb{E}_{x \sim \mathcal{P}} \left[\frac{\mathcal{A}_2(x)}{OPT(x)} \right] \geq \frac{1}{2}(\sqrt{n} + 1).$$

Finally, this concludes the proof also for **Group 2**. It's easy to see that any algorithm is either $\mathcal{A}_1, \mathcal{A}_2$, inside **Group 1** or inside **Group 2**, hence we have concluded the proof. \square

Now, the result simply follows from Yao's principle,

$$\max_{x \in \mathcal{X}} \frac{\mathbb{E}_{a \sim \mathcal{A}}[a(x)]}{OPT(x)} \stackrel{(i)}{=} \max_{x \in \mathcal{X}} \mathbb{E}_{a \sim \mathcal{A}} \left[\frac{a(x)}{OPT(x)} \right] \stackrel{(ii)}{\geq} \min_{\mathcal{A} \in \text{det. algos}} \mathbb{E}_{x \sim \mathcal{P}} \left[\frac{\mathcal{A}(x)}{OPT(x)} \right] \stackrel{(iii)}{\geq} \frac{1}{2}(\sqrt{n} + 1) \in \Omega(\sqrt{n})$$

Where (i) follows from linearity of expectation, (ii) is a direct application of Yao's principle using $c(a, x) = \frac{a(x)}{OPT(x)}$. (iii) follows from Lemma 1. Being the first term of the last equation the competitive ratio of an arbitrary randomized algorithm \mathcal{A} , we have showed that it is in $\Omega(\sqrt{n})$, and therefore, we have concluded the proof.